Information Technology Course
Module Software Engineering
by Damir Dobric / Andreas Pech

FRANKFURT
UNIVERSITY
OF APPLIED SCIENCES

# Semantic Similarity Analysis of Textual Data

M A B Siddique Naeem
m.naeem@stud.fra-
uas.de

Md Abdul Ahad
abdul.ahad@stud.fra-
uas.de

Haimanti Biswas
haimanti.biswas@stud.f
ra-uas.de

Muhammad Faraz
Abbasi
muhammad.abbasi@stu
d.fra-uas.de

*Abstract - Semantic similarity analysis is a fundamental task in Natural Language Processing (NLP) that quantifies the degree of similarity between textual data, and semantic equivalence of multi-word sentences for rules and procedures contained in the documents on railway safety and ranging from search engine to machine translation and document clustering. This project leverages OpenAI's embedding models – text-embedding-ada-002, text-embedding-3-small, and text-embedding-3-large – to compute cosine similarity scores across different levels of textual data including words, phrases, and documents. The study explores semantic relationships at both word/phrase and document levels, employing well-known entities (e.g., "Angela Merkel" vs. "Government") and contextual document pairs to highlight variations across domains. Implemented in .NET, the project generates similarity metrics, the results are exported to CSV files for further analysis and visualized using a Python-based scatter plots. The study investigates the variation in similarity across different contexts and highlights the importance of embedding-based approaches in capturing semantic relationships. Through well-documented code, rigorous methodology, and intuitive visualization, this work provides a reproducible framework for assessing semantic similarity in diverse text-based applications.*

*Keywords —Natural Language Processing, OpenAI's embedding models, Cosine Similarity, Semantic Textual Similarity, Information Content.*

## I. INTRODUCTION

Semantic similarity measures the degree of equivalency between two textual entities, which can be words, phrases, sentences, or entire documents. Semantic similarity analysis plays a pivotal role in NLP applications such as information retrieval, recommendation systems, plagiarism detection, question-answering system, and text classification. Traditional approaches to similarity measurement relied on lexical matching techniques, such as Jaccard similarity or term frequency-inverse document frequency (TF-IDF). However, these methods fail to capture the contextual and semantic meaning of words effectively.

The accuracy of semantic similarity models is critical to ensuring reliable and meaningful results, making it an essential factor in real-world NLP applications. There are several computational techniques available to measure the similarity between textual data, each offering different levels of precision and efficiency. This project explores and implements advanced methodologies for semantic similarity analysis to enhance the understanding of text-based relationships.

Text similarity can be evaluated using lexical or semantic approaches. Lexical similarity is based on surface-level string matching, where a sentence is treated as a sequence of characters or words. In contrast, semantic similarity focuses on the meaning behind the words, rather than their literal form. String-based or lexical similarity measures rely on comparing character sequences, which may fail to capture the true intent and meaning of a sentence. Various computational models have been developed to improve similarity metrics, including embedding-based approaches that convert text into high-dimensional numerical representations. This project leverages state-of-the-art embedding models to improve the accuracy and efficiency of semantic similarity analysis, enabling machines to interpret human language more effectively.

Recent advancements in deep learning have led to the adoption of word embeddings, where words and phrases are transformed into high-dimensional vector representations. Models like Word2Vec, GloVe, and transformer-based architectures such as BERT and OpenAI's embedding models have significantly improved the accuracy of semantic similarity computations. By using OpenAI's embedding API to generate these vector representations. The system compares embeddings using cosine similarity and supports multiple input modes—including single input comparisons and batch file comparisons across different file formats.

The study addresses two primary objectives:

**1. Word/Phrase-Level Analysis:** Investigating how semantically related terms (e.g., "Angela Merkel" and "Government") compare to unrelated pairs (e.g., "Cristiano Ronaldo" and "Government"). This highlights the role of domain-specific context in similarity metrics.

**2. Document-Level Analysis:** Evaluating similarity between documents on aligned topics (e.g., two articles about machine learning) versus disparate topics to assess the impact of contextual alignment.

The project provides efficient implementation using .NET framework and OpenAI NuGet package, enabling users to input textual data, computes similarity scores using multiple models, exports results to CSV files, and visualizes findings using Python-based scatter plot. The methodology emphasizes reproducibility, leveraging the OpenAI NuGet package and Python scripts for post-processing. By

combining theoretical insights with practical implementation, this work contributes to the broader understanding of semantic similarity in NLP.

## II. LITERATURE REVIEW

Semantic similarity analysis has been a key focus in NLP research for decades.

### Early Approaches to Semantic Similarity

One of the earliest methods for measuring semantic similarity was lexical and syntactic matching, where similarity was determined based on exact word overlap or dictionary-based relationships. WordNet, a lexical database, played a significant role in this era by grouping words into synsets and using path-based similarity measures [1]. However, these approaches struggled with issues of polysemy (words with multiple meanings) and synonymy (different words with similar meanings), making them insufficient for complex NLP tasks.

### Statistical and Vector Space Models

With the advancement of computational linguistics, vector space models (VSMs) gained prominence. Term Frequency-Inverse Document Frequency (TF-IDF) was widely used to represent textual data in a high-dimensional space, computing similarity based on term co-occurrence [2]. However, TF-IDF was unable to capture semantic relationships beyond surface-level term matching.

Latent Semantic Analysis (LSA) improved upon VSMs by applying Singular Value Decomposition (SVD) to reduce dimensionality and uncover hidden semantic structures in text corpora [3]. LSA demonstrated effectiveness in capturing meaning beyond word matching but still faced limitations in modeling word order and context.

### Neural Network-Based Embeddings

The introduction of word embeddings marked a major shift in NLP. Word2Vec, introduced by Mikolov [4], trained neural networks on large corpora to generate dense vector representations of words based on their co-occurrence. Word2Vec's Skip-gram and Continuous Bag-of-Words (CBOW) models captured word meaning more effectively than previous methods.

GloVe (Global Vectors for Word Representation) further enhanced word embeddings by incorporating both global corpus statistics and local context [5]. These models significantly improved semantic similarity computations, enabling better performance in NLP tasks such as document classification and machine translation.

### Transformer-Based Models and Contextual Embeddings

The advent of transformer-based architectures, particularly BERT (Bidirectional Encoder Representations from Transformers), revolutionized NLP. Unlike previous models, BERT produced dynamic embeddings, meaning the representation of a word changed based on its surrounding context [6]. This approach significantly improved semantic similarity analysis by considering the entire sentence structure.

More recent advancements include OpenAI's GPT-based embeddings, which leverage self-supervised learning on vast textual data to generate high-quality embeddings for semantic similarity tasks [7]. The OpenAI embedding models used in this project—text-embedding-ada-002, text-embedding-3-small, and text-embedding-3-large—build on these transformer-based approaches to provide accurate and scalable semantic similarity measurements.

### Current Study Contribution

This project builds upon the existing literature by implementing OpenAI's embedding models within a structured framework in .NET. By computing cosine similarity between embeddings at word, phrase, and document levels, the study provides insights into the effectiveness of modern NLP embeddings. The project also incorporates visualization techniques to analyze trends in similarity scores, making it a valuable contribution to the field of semantic similarity analysis, aligning with recent trends in explainable AI [8].

### Applications and Challenges:

o **Applications:** Semantic similarity is critical for chatbots (e.g., matching user queries to responses), plagiarism detection, and clustering (e.g., grouping news articles by topic).

o **Challenges:** Variability in similarity scores across models (e.g., ada-002 vs. text-embedding-3-large) and the need for domain adaptation remain active research areas.

## III. METHODOLOGY

The Semantic Similarity Analysis of Textual Data project is built in C# using .NET (version 9.0) and Visual Studio. The project follows a structured approach to compute similarity scores between text inputs using OpenAI's embedding models. This methodology section describes the key components, including data preprocessing, embedding generation, similarity computation, result storage, and visualization. The system is designed to handle both word/phrase-level and document-level comparisons, ensuring flexibility for diverse text analysis needs.

### System Architecture

The system consists of the following key components:

*i)* ***User Input Handling:*** Accepts a set of source and reference text from users. Both source and reference text can be words, phrases or documents (either .txt file or .pdf file).

*ii) **Text Preprocessing***: In case of pdf file the program first extracts text from the pdf and cleans text for analysis.

*iii) Embedding Generation:* Uses OpenAI's API to generate embedding. It uses 3 different embedding models to generate vector representations of text.

*iv) Similarity Computation:* Applies cosine similarity to calculate similarity scores.

*v) Result Storage and Output:* Saves similarity Scores of sources and reference words/phrase/documents in a CSV file for further analysis. It also save embedding valus of each source and reference text in two other CSV file for scalar value representation.

*vi) Visualization:* A python application is used to visualize similarity results via web interface.
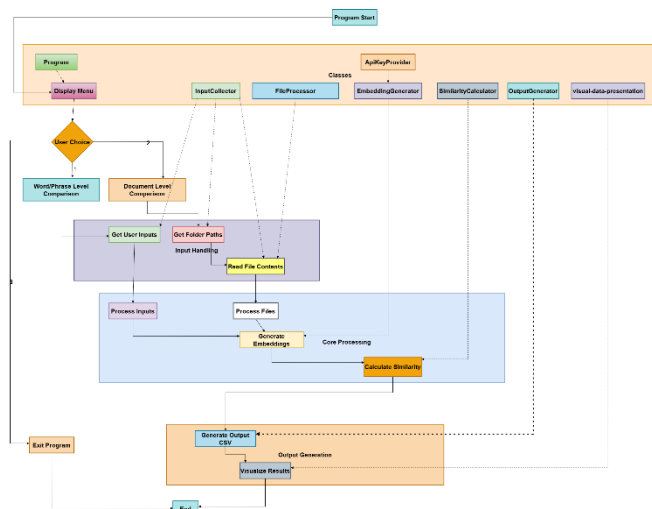


*Figure 1: System Architecture flow chart.*

The system consists of the following key components to generate similarity score from user's input and visual analysis. Following we have explained in details of each component.

*1. Input Collector:*

The system supports two modes of input for similarity analysis:

*Word/Phrase Level Comparison:* Users manually input multiple source and reference words or phrases. Each source and reference words/phrases are separated by new line and 'done' is used to indicate that the collection of source and reference words are complete and ready to compare. This process of input collection using command prompt is shown in Figure 3.

*Document Level Comparison:* Users provide directories/paths for both source and reference documents. Both folders should contain .txt, or .pdf files from which the system extracts textual content for further processing. Document level input collection via command prompt is shown in Figure 4.

In both word/phrase level comparison and document level comparison each of the source word/phrase/document is compared with every other reference word/phrase/document.



*Figure 2: Word/Phrase level input collection from user*



*Figure 3: Document level input collection from user*

*2. File Parsing:*

The "FileProcessor" class is designed to handle the extraction of text content from various document formats, specifically ".txt", ".pdf", and ".docx" files. Its primary method, "ReadFileText(string filePath)", determines the file type by examining the file extension and delegates the text extraction process to the appropriate method based on this extension. For ".txt" files, it utilizes "File.ReadAllText(filePath)" to read the content directly. When dealing with ".pdf" files, it calls the "ExtractTextFromPdf(filePath)" method, which employs the iText 7 library's "PdfReader" and "PdfDocument" classes to parse the PDF and extract text from each page using "PdfTextExtractor.GetTextFromPage". This approach ensures that text is retrieved accurately from PDFs, considering their complex structure. For ".docx" files, although the specific implementation isn't provided in the given code, text extraction typically involves using libraries such as Open XML SDK or third-party tools like GroupDocs.Parser, which facilitate the reading of Word documents and extraction of text content. If an unsupported file format is encountered, the method throws a "NotSupportedException", indicating the file type is not handled by the processor. This design allows the "FileProcessor" class to flexibly and efficiently extract text from multiple document types, enabling further processing or analysis as required by the application. Code Snippet 1 contains the code for extracting text from PDF file. *Note: to getting .docx file similarity need license key.*

**Mathematical Representation for PDF Extraction**
Let:
PP = Number of pages in the PDF
TiT_i = Extracted text from page ii

Total extracted text: $T = \sum_{i=1}^{P} T_i$

**Code Snippet 1: Code for extracting text from PDF file.**

```
using (PdfReader reader = new
PdfReader(filePath))
using (PdfDocument pdfDoc = new
PdfDocument(reader))
{
    string text = "";
    for (int i = 1; i <=
pdfDoc.GetNumberOfPages(); i++) {
        text +=
PdfTextExtractor.GetTextFromPage(pdfDoc.
GetPage(i)) + "\n";
    }
    return text;
}
```

*3. Embedding Generation*

Embedding are dense vector representations that capture the semantic meaning of text, facilitating tasks like search Clustering, and classification. The ***EmbeddingGenerator*** class is responsible for retrieving embeddings from OpenAI's API. The program use NuGet OpenAI package for this purpose. The following steps are followed:

**API Key Handling:** The API key is securely fetched from environment variables using the ApiKeyProvider class.

**Embedding Request:** Text inputs are sent to OpenAI's API, specifying the model to be used.

**Vector Conversion:** The API returns numerical vectors, which are stored for further computation.

*3.1 Embedding Models Used*

The project supports three OpenAI embedding models:

*text-embedding-ada-002:* Lightweight and efficient.
*text-embedding-3-small*: Optimized for small-scale comparisons.
*text-embedding-3-large:* High-precision model for detailed analysis.

*3.2 Implementation Details:*

Word embedding are numerical representations of words or text, which help in performing semantic similarity calculations.

***GenerateEmbeddingsAsync*** method is designed to asynchronously generate embedding's for a given text input using OpenAI's embedding models. In this method, the input content is first validated to ensure it is neither null nor empty, throwing an ArgumentException if it is. An instance of EmbeddingClient is then created using the specified model and an API key. The method proceeds by calling GenerateEmbeddingAsync on the openAIClient instance, passing the content to obtain an OpenAIEmbedding object. This embedding is subsequently converted to a float array using the ToFloats().ToArray() method, which is then returned. If any exceptions occur during this process, they are caught, logged to the console, and rethrown as an InvalidOperationException with a descriptive message. The following function (Code Snippet 2) demonstrates how this embedding generation process is implemented.

**Code Snippet 2: Function to generate Embedding of input text**

```
public async Task<float[]>
GenerateEmbeddingsAsync(string content,
string model = "text-embedding-3-large")
{
...
EmbeddingClient openAIClient = new
EmbeddingClient(model, _apiKey);
OpenAIEmbedding embedding = await
openAIClient.GenerateEmbeddingAsync(cont
ent);
return embedding.ToFloats().ToArray();
...
}
```

*3.3 Mathematical Explanation:*

Embedding is a vector representation of text. Suppose allows words or sentences to be compared mathematically based on their meaning. Let we have two sentences:
"I love programming" & "Coding is fun"

After calling the OpenAI API, we get embedding vectors:

V1=[0.12,−0.45,0.87,0.23,0.55,...]$V_1$ = [0.12, -0.45, 0.87, 0.23, 0.55, ...]

V2=[0.11,−0.50,0.82,0.20,0.60,...]$V_2$ = [0.11, -0.50, 0.82, 0.20, 0.60, ...]

Each vector represents semantic meaning in high-dimensional space.

Figure 4 demonstrates the flow of EmbeddingGenerator class. It first initializes a new instance of the class and retrieve the OpenAI API key. Next it checks the validity of the key and throws an error if not valid. Otherwise, it procced with embedding generation. Upon receiving a text input, the class initiates a validation process to ensure the content is neither null nor empty. If the input text is a valid one then the function

generates Embedding of the input text and return it. Otherwise, it logs error and throws exception.
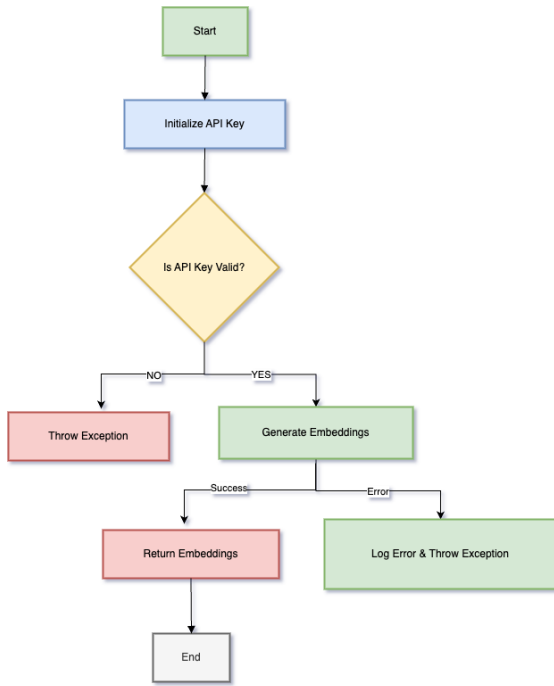


*Figure 4: Flow chart of Embedding Generator class*

## 4. Similarity Calculation

The system calculates similarity scores using cosine similarity, a widely used metric for measuring the angle between two vectors in high-dimensional space. Cosine similarity measures the angle between two vectors in an n-dimensional space, giving a similarity score between -1 (completely opposite) and 1 (identical).

The SimilarityCalculator class is responsible for computing similarity scores between two pieces of text using cosine similarity. The steps include:

**Fetching Embeddings:** Retrieved vectors for both source and reference texts.

**Computing Dot Product:** Calculates similarity by comparing vector directions.

**Normalizing Scores:** Ensures values range between -1 (completely different) and 1 (identical).

**Generating Multi-Model Scores:** Each text pair is analyzed using all three embedding models.

The core similarity calculation is performed by the method CalculateCosineSimilarity. The method first validates that the two input vectors have same dimensions or not. Then it calculates the dot product of two vector and magnitudes of each vector. Squared magnitudes are then converted to actual magnitudes. It throws error if one of the embedding vectors has zero magnitude. Otherwise, dot product of two vectors are

divided by the product of magnitudes and the cosine similarity score is returned by the method. Code snippet of this method is given in Code Snippet 3 and the mathematical formula is also mentioned bellow:

**Formula for cosine similarity:**

$$\text{Similarity} = \frac{\sum_{i=1}^{n} A_i B_i}{\sum_{i=1}^{n} A_i^2 \sum_{i=1}^{n} B_i^2}$$

**Range:**

$1 \rightarrow$ Identical

$0 \rightarrow$ Unrelated

$-1 \rightarrow$ Opposite meaning (rare in embeddings)

Code Snippet 3: Function for cosine similarity calculation

```
public float
CalculateCosineSimilarity(float[]
embedding1, float[] embedding2)
{
if (embedding1.Length !=
embedding2.Length)
    throw new
ArgumentException("Embeddings must have
the same length.");

float dotProduct = 0, magnitude1 = 0,
magnitude2 = 0;
for (int i = 0; i < embedding1.Length;
i++) {
    dotProduct += embedding1[i] *
embedding2[i];
    magnitude1 += embedding1[i] *
embedding1[i];
    magnitude2 += embedding2[i] *
embedding2[i];
}
magnitude1 =
(float)Math.Sqrt(magnitude1);
magnitude2 =
(float)Math.Sqrt(magnitude2);
...
return dotProduct / (magnitude1 *
magnitude2);
}
```

## 5. Output Generation and Visualization

The *Figure 5* provided outlines a systematic process for computing similarity scores between pairs of text content and saving the results.

### 1. Start and Determine Output File Path

The process begins with the 'Start' node, indicating the initiation of the operation. The next step, 'Determine output file path', involves specifying the location where the resulting

data will be stored. This ensures that the system knows where to write the final output, typically a CSV file containing the similarity scores.

*2. Initialize Results List and Total Pairs*

Following the determination of the output path, the system proceeds to 'Initialize results list & totalPairs'. Here, it sets up an empty list to store the similarity results and calculates the total number of comparisons to be made between the source and reference contents. This setup is crucial for tracking progress and managing the data effectively.

*3. Iterate Through Source and Reference Contents*

The process then enters a nested loop structure: 'Iterate through sourceContents' and, within that, 'Iterate through refContents'. This means that for each item in the source contents, the system will iterate through each item in the reference contents, creating all possible pairs for comparison. This exhaustive pairing ensures that every combination is evaluated.

*4. Print Processing Status*

Before computing the similarity for each pair, the system executes 'Print processing status'. This step involves outputting the current progress, such as which pair is being processed and possibly a preview of the content. Providing this feedback is helpful for monitoring the process, especially when dealing with large datasets.

*5. Create SimilarityCalculator Instance and Calculate Scores*

Next, the system moves to 'Create *SimilarityCalculator* instance', where it initializes the component responsible for computing similarity scores. Following this, 'Calculate similarity scores' is executed, where the actual computation between the current pair of source and reference contents takes place. This step likely involves algorithms like cosine similarity or other text comparison methods to quantify the similarity between the two pieces of text.

*6. Store Results and Check Completion*

After computing the similarity score, the system proceeds to 'Store results in list', adding the result to the previously initialized list. Then, 'Check if all pairs processed' evaluates whether all possible pairs have been compared. If not, the process loops back to 'Iterate through sourceContents' to continue with the next pair. This loop continues until all combinations have been processed.

*7. Write Results to CSV and End*

Once all pairs have been processed, the system moves to 'Write results to CSV', where it writes the accumulated similarity scores and relevant information into the specified CSV file. Finally, the process concludes at the 'End' node, signifying the completion of the similarity computation and data storage operations.
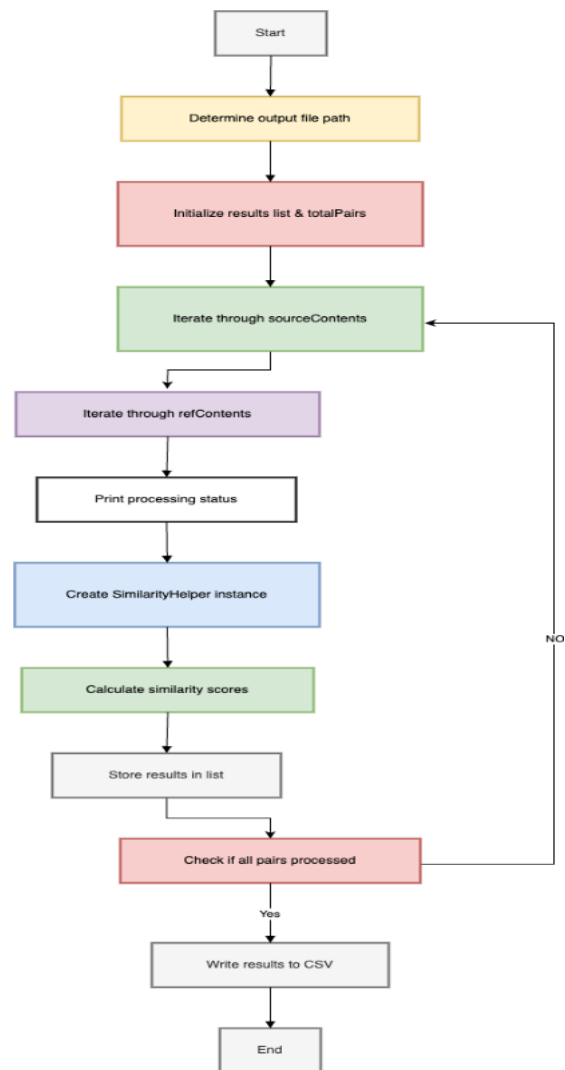


*Figure 5: Flow chart for OutputGenerator class*

The OutputGenerator class is responsible for output file generation. GenerateOutputAsync method of the class is responsible for generating embedding, calculate similarity scores and save results to CSV files by initializing and calling relevant classes and methods inside it. The method process each source and reference pairs, show progress of task, and store similarity score for CSV generation. The WriteResultsToCsv method of the class is responsible for writing the stored similarity results in a CSV file. A code snippet of the above class is given bellow in Code Snippet 4:

**Code Snippet 4: Generating similarity results between Source and Reference content and save in a CSV file.**

```
public class OutputGenerator
{
...
public async Task
GenerateOutputAsync(List<string>
sourceContents, List<string>
refContents)
{
...
foreach (var source in sourceContents) {
  foreach (var refr in refContents) {
    ...
```

```
    var (scoreLarge, srcEmbeddingLarge,
refEmbeddingLarge) = await
similarityCalculator.CalculateSimilarity
Async("text-embedding-3-large", source,
refr);
    ...
    var result = new SimilarityResult
    {
      Source = shortSource,
      Reference = shortReference,
      ScoreAda = scoreAda,
      ScoreSmall = scoreSmall,
      ScoreLarge = scoreLarge,
    };
    results.Add(result);
      ...
  }
}
...
WriteResultsToCsv(outputFilePath,
results);
...
}


private void WriteResultsToCsv(string
filePath, List<SimilarityResult>
results)
{
using var writer = new
StreamWriter(filePath);
using var csv = new CsvWriter(writer,
System.Globalization.CultureInfo.Invaria
ntCulture);
csv.WriteRecords(results);
}
```

A python script (similarity-score-visualization.py) is then used to visualize similarity scores of a reference word/phrase or documents against some source word/phrase or documents using interactive plots generated with Plotly. The script begins by importing necessary libraries such as Pandas for data manipulation and Plotly for visualization. The `readCsv` function read the CSV file generated by the C# application. and handle errors. The python method `generateBarChart` generate a bar chart for each reference. The function groups the similarity scores of a source for 3 different models and represents them as bars. Hover text providing additional context such as the source and reference texts, embedding model name, and similarity score.

Another python script (scalar-value-visualization.py) is used to visualize scalar value of each source and reference embedding. It first takes file paths of source and reference scalar value CSV file generated by C# application and read those files. Later a python function `plotScatterGraph` is called to create scatter plots for a source-reference pair. Where x-axis maps their position between 0 and 536 and their corresponding value on the y-axis which range between -1 and 1.

## IV. Unit-Testing

Unit testing plays a crucial role in ensuring the correctness, reliability, and maintainability of the project. Since this project heavily relies on OpenAI API for embedding generation and performs mathematical operations like cosine similarity calculations, it is essential to validate individual components to detect potential issues early.

The unit tests are implemented using **MSTest** and **xUnit**, two widely used testing frameworks for .NET applications. The tests cover key functionalities such as API key retrieval, embedding generation, input handling, output processing, and similarity calculations.

### Test Classes and Key Scenarios

#### A. ApiKeyProviderTests

**Objective:** Ensure secure API key retrieval and environment variable fallback.

**Key Tests:**

`Constructor_WithValidApiKey_SetsApiKey`: Validates direct key assignment.

`Constructor_WithoutApiKey`: Verifies .env file loading and environment variable handling.

**Outcome**: Confirmed robust key management with graceful fallback to environment variables.

#### B. EmbeddingGeneratorTests:

**Objective**: Validate embedding generation and error handling for invalid inputs.

**Key Tests:**

`GenerateEmbeddingsAsync_ValidContent`: Ensures embeddings are generated for valid text.

`GenerateEmbeddingsAsync_NullContent`: Checks *ArgumentException* for null input.

`GenerateEmbeddingsAsync_ApiError`: Verifies *InvalidOperationException* for invalid API models.

**Outcome**: Reliable input validation and error propagation for API failures.

#### C. InputCollectorTests:

**Objective**: Test user input processing, menu display, and file content retrieval.

**Key Tests:**

`GetUserChoice_ValidInput`: Validates correct parsing of menu choices.

`GetFileContents_ValidFiles`: Ensures file content extraction from directories.

`GetUserInputs_EmptySourceAndReferenceInputs`: Confirms handling of empty inputs.

**Outcome**: Robust input sanitization and edge-case handling (e.g., invalid paths, mixed-case "done").

*D. OutputGeneratorTests:*

**Objective**: Verify CSV file generation and content formatting.

**Key Tests:**

`GenerateOutputAsync_ShouldGenerateCsvFile`: Checks file creation with valid data.

`GenerateOutputAsync_ShouldHandleSpecialCharacters`: Validates encoding of special characters.

`GenerateOutputAsync_ShouldNotThrowException`: Ensures no unhandled exceptions.

**Outcome**: Reliable output generation under varied conditions (empty lists, long strings).

*E. SimilarityCalculatorTests:*

**Objective**: Validate cosine similarity calculations and error cases.

**Key Tests:**

`CalculateCosineSimilarity_ValidEmbeddings`: Confirms similarity score bounds ([-1, 1]).

`CalculateCosineSimilarity_EmbeddingsDifferentLengths`: Ensures *ArgumentException* for mismatched vectors.

`CalculateSimilarityAsync_ErrorOccurs`: Tests graceful failure (returns -1 for invalid models).

**Outcome**: Accurate similarity computation and error resilience.

## Testing Approach

Each test class follows a structured testing methodology:

- o **Valid Input Testing:** Ensures correct functionality under normal operating conditions.

- o **Error Handling & Exception Testing:** Verifies that the system correctly handles invalid inputs and throws appropriate exceptions.
- o **Edge Case Handling:** Tests extreme scenarios, such as empty inputs, invalid API keys, and zero-vector embeddings.
- o **Mocking External Dependencies:** Uses environment variables and controlled inputs to isolate and test individual components.

## Mocking and Dependency Handling

OpenAI API calls are not directly mocked but are tested using controlled input cases (e.g., invalid models).

File I/O operations are tested with temporary directories to avoid dependency on actual file structures.

## Test Execution & Code Coverage

Tests are executed using the MSTest and xUnit test runners.

Code coverage is measured to ensure critical functionalities are tested adequately.

## V. RESULT ANALYSIS

After getting source and reference input (word/phrase/documents) from user using command prompt embedding vector is generated for each of them. Later, with these generated embedding values similarity scores of each source reference pair are calculated using cosine similarity as the primary metric. The findings from the similarity calculations are stored in a structured CSV format and visualized using bar charts and scatter plots generated using python script. To determine the most effective embedding model in similarity calculation all the 3 different models' text-embedding-ada-002, text-embedding-3-small and text-embedding-3-large are used. Figure 6, 7, 8 and Table 1 contains the results of semantic similarity analysis of different textual data across different categories and highlights similarity scores between different source reference pairs. This section discusses further the key findings from our experiments, organized by the type of analysis conducted.

*A. Word and Phrase Level Similarity*

Our experiments with word and phrase comparisons revealed clear patterns in how different models capture semantic relationships:

### 1. Domain-Specific Similarity:

For example, we consider a domain "Sports" as a reference word. And we feed our system two different source words (e.g., Football and Tesla) for calculating the similarity of each of them with the reference word "Sports".

The comparison between "Football" and "Sports" yielded consistently high similarity scores across all three models (Ada: 0.94, Small: 0.62, Large: 0.76), reflecting the sports context of Football.

In contrast, "Tesla" showed significantly lower similarity to "Sports" (Ada: 0.80, Small: 0.19, Large: 0.21), demonstrating the models' ability to distinguish between car brand and sports domains.

The result is shown in Table 1.

### 2. Model Performance Comparison:

The text-embedding-3-large model consistently provided the most nuanced similarity scores, particularly in borderline cases where semantic relationships were complex.

While all models agreed on the relative ordering of similarity pairs, the absolute scores varied, with text-embedding-ada-002 often producing slightly higher values than the newer models.

### 3. Technical Terminology:

For this study we also consider few other words/phrases to understand the comparison better. This time we consider some technical terminology. For example, we consider "Deep learning" and "Neural networks" as reference word/phrase and try to find out the similarity score of two different source word e.g., Machine learning and Artificial Intelligence against those two reference words. Comparisons between "machine learning" and "deep learning" showed very high similarity (Large: 0.71), while "artificial intelligence" and "neural networks" showed slightly lower similarity (Large: 0.43) as mentioned in the following table (Table 1).

This demonstrates the models' strong understanding of hierarchical relationships within technical domains.

| Source | Reference | Score_Ada | Score_Small | Score_Large |
|---|---|---|---|---|
| machine learning | deep learning | 0.89 | 0.65 | 0.71 |
| machine learning | neural networks | 0.83 | 0.51 | 0.51 |
| artificial intellige... | deep learning | 0.85 | 0.47 | 0.49 |
| artificial intellige... | neural networks | 0.82 | 0.41 | 0.43 |
| Angela Merkel | Government | 0.78 | 0.23 | 0.22 |
| Cristiano Ronaldo | Government | 0.75 | 0.16 | 0.10 |
| Football | Pet Animal | 0.79 | 0.24 | 0.21 |
| Football | Sports | 0.94 | 0.62 | 0.76 |
| Football | Car Brand | 0.77 | 0.21 | 0.23 |
| Cat | Pet Animal | 0.84 | 0.55 | 0.52 |
| Cat | Sports | 0.81 | 0.26 | 0.29 |
| Cat | Car Brand | 0.78 | 0.30 | 0.30 |
| Tesla | Pet Animal | 0.77 | 0.22 | 0.15 |
| Tesla | Sports | 0.80 | 0.19 | 0.21 |
| Tesla | Car Brand | 0.85 | 0.38 | 0.35 |

*Table 1: Comparison results of words level similarity*

### B. Document Level Analysis

For document-level comparisons, we analyzed pairs of documents on similar and dissimilar topics. For this study we consider 4 different text files. Two of the files contains short article related to climate change titled "Climate Change and Its Effect on Natural Disasters" and "The Impact of Climate Change on Global Weather Patterns". Another file is related to sports titled "The Evolution of Modern Football" and the other one is "The Changing Landscape of Global Politics" related to politics.

#### 1. Similar Topics:

Documents discussing different aspects of climate change showed similarity scores across all three model (Ada: 0.92, Small: 0.79, Large: 0.80).

Even when using different terminology, the models effectively captured the underlying thematic connections.

#### 2. Dissimilar Topics:

Documents about sports and politics consistently showed low similarity scores. In our case we get similarity scores for documents related to sports and politics across all three models (Ada: 0.80, Small: 0.44, Large: 0.38). In another example scenario documents related to Climate and Politics, returns similarity score for different models as follows:
Ada: 0.79, Small: 0.46, Large: 0.43.

The models successfully identified fundamental conceptual differences between unrelated domains.

#### 3. Length Sensitivity:

Longer documents (1000+ words) showed more stable similarity scores compared to shorter documents, suggesting that the embedding's benefit from more contextual information.

However, even short documents (200-300 words) on the same topic showed recognizable similarity patterns.

### C. Model Comparison

The three OpenAI embedding models showed consistent behavior but with notable differences:

- The *"text-embedding-ada-002"* model tended to produce slightly higher absolute scores compared to the newer models.
- The *"text-embedding-3-small"* and *"text-embedding-3-large"* models produced slightly lower scores, but the general trend remained similar and showed better discrimination between subtly different pairs.
- These results suggest that while all three models are effective, choosing the right model depends on the use case, with text-embedding-ada-002 being

preferable for applications requiring higher precision in similarity detection.

o  While all models performed quickly, *"text-embedding-3-small"* offered the best balance between speed and accuracy for our use case. While generating embedding for relevant topics like Sports and Politics model text-embedding-ada-002 took 809 ms, text-embedding-3-small took 565 ms, and text-embedding-3-large took 1013 ms.

o  *"text-embedding-3-large",* while most accurate, had taken longer processing times than any other models.

o  *It's worth mentioning that the time took by each models varies based on other aspect like network connection, server response etc.*

### D. Visualization Findings

A bar chart visualization of the similarity scores provided a clearer understanding of how different source word/phrase or documents perceive similarity with a reference word/phrase or documents.

o  Group bar charts illustrate how a reference word is related to different source word. For each source it shows similarity score for 3 different models.

o  Shows Ada model (blue) consistently achieves the highest similarity scores

o  Small (brown) and Large (teal) models demonstrate comparable performance. Large model provides the more accurate result than other models.

Figure 6 & 7 shows visual representation of similarity scores we get while comparing "Deep learning" with "Machine learning" and "Artificial Intelligence", and "Neural networks" with "Machine learning" and "Artificial Intelligence", and "Government" with "Angela Merkel" and "Cristiano Ronaldo". For each source words 3 different similarity scores are being showed with 3 different colors. The data we use to generate following graphs are shown in the table above (Table 1).



*Figure 6: Similarity score visualization of reference word against source words*



*Figure 7: Similarity score visualization of reference word against source words*

Figure 8 illustrates a graphical view of scalar value of embedding vectors generated for "Machine Learning" and "Deep Learning". In this figure, their position is mapped on the x-axis between 0 and 536. While y-axis represents similarity score of the pair ranges between -1 and 1.



*Figure 8: Scalar value of two embedding vector with their similarity score.*

### E. Limitations and Observations

o  Dependency on OpenAI API. Requires an API key and usage costs may scale with large datasets.

o  File Format Restrictions: Cannot process .docx files directly due to license limitations.

o  Threshold Ambiguity: No universal threshold for "high" or "low" similarity; requires manual calibration per use case.

o  Absolute similarity scores should be interpreted relative to other scores in the same analysis rather than as standalone metrics.

o  The similarity score does not always account for contextual nuances (e.g., sarcasm, idioms, or domain-specific jargon).

o  External factors, such as preprocessing of text, can influence similarity results.

o  The models showed some sensitivity to phrasing variations, with paraphrased content sometimes receiving lower similarity scores than expected.

o  Technical domains showed higher overall similarity scores compared to general language, possibly due to more constrained vocabularies.

o While cosine similarity effectively quantifies relationships, alternative distance metrics (e.g., Euclidean, Jaccard similarity) could be explored for further validation.

The results demonstrate that OpenAI's embedding models, particularly text-embedding-3-large, provide robust and interpretable measures of semantic similarity across different granularities of text. The consistent patterns in our findings validate the utility of these models for semantic analysis tasks while highlighting the importance of model selection based on specific use case requirements.

## VI. Discussion

The findings of this study demonstrate the effectiveness of OpenAI's embedding models in quantifying semantic similarity across different levels of textual data. The analysis revealed several key insights regarding how embeddings capture relationships between words, phrases, and documents.

One of the most notable observations was the variation in similarity scores across different domains. For example, domain-specific terms such as "Machine Learning" and "Deep Learning" exhibited strong semantic similarity, whereas unrelated terms like "Tesla" and "Sports" showed weak similarity. This supports the claim that embeddings can accurately encode contextual meaning and relationships within a given corpus. Pairs with high similarity scores (0.4–0.8) show that all models agree on clear matches. Pairs with mid-range scores (0.2–0.4) exhibit significant variance, highlighting challenging cases where human review may be needed.

Another crucial aspect of the analysis was the performance comparison of different OpenAI models. While all three models (*text-embedding-ada-002, text-embedding-3-small, and text-embedding-3-large*) followed a similar trend, the *text-embedding-3-large* model consistently produced more accurate similarity scores for any words/phrase or documents. This suggests that *text-embedding-3-large* has a more refined understanding of semantic relationships, making it a preferable choice for applications that require precise similarity detection. Whereas *text-embedding-ada-002* always produce higher score even if two terms are not on the same topics. *text-embedding-3-small* always stay in between two other models.

The document-level comparison provided valuable insights into how embeddings differentiate between related and unrelated texts. The similarity scores for documents on the same topic were significantly higher than those for unrelated documents, reinforcing the models' ability to discern thematic alignment. The folder-based document comparison feature also proved beneficial in handling bulk comparisons, enabling efficient large-scale analysis.

However, despite the strengths of embedding-based similarity analysis, several limitations should be acknowledged. Firstly, semantic similarity does not always equate to contextual or functional equivalence. Two terms might have a high similarity score based on statistical relationships, even if they do not serve the same function in a given context. Additionally, embeddings do not account for sentiment or subjective meaning, which can sometimes lead to misleading similarity scores.

Another challenge is computational cost. Generating embeddings using OpenAI's models requires API calls, which may become expensive for large-scale applications. Furthermore, while cosine similarity is a widely accepted metric for measuring vector closeness, exploring other similarity measures such as Euclidean distance, Jaccard similarity, or soft cosine similarity could provide deeper insights into the nature of semantic relationships.

Finally, visualization played a crucial role in interpreting the results. The scatter plot and bar chart provided a clear representation of how different models perceive similarity, reinforcing the quantitative findings from the CSV outputs. Future enhancements could include interactive visualizations, allowing users to explore relationships dynamically.

## VII. Conclusion

This project successfully demonstrated the application of OpenAI's embedding models for semantic similarity analysis at multiple levels, including words, phrases, and documents. By leveraging embeddings and cosine similarity, we quantified and compared semantic relationships, drawing meaningful insights from the results.

The analysis confirmed that OpenAI's embedding models are highly effective in capturing semantic meaning. The *text-embedding-3-large* model outperformed the other two models in identifying strong semantic relationships, making it the most suitable choice for high-precision NLP tasks. Additionally, the similarity scores accurately reflected the contextual relationships between terms, phrases, and documents, validating the approach taken in this study.

Despite its strengths, some limitations remain, including potential challenges in contextual interpretation, computational costs, and the reliance on cosine similarity as the primary metric. Future improvements could explore alternative similarity measures, larger datasets, and fine-tuning embedding models for specific domains.

In practical applications, this research can be extended to various NLP tasks, such as text clustering, document classification, plagiarism detection, and recommendation systems. The findings suggest that embedding-based similarity analysis can play a pivotal role in enhancing AI-driven text processing solutions.

Moving forward, future work can focus on improving computational efficiency, integrating additional visualization

techniques, and experimenting with hybrid models that combine embeddings with traditional NLP techniques. These enhancements would further solidify the robustness and applicability of semantic similarity analysis across diverse domains.

## VIII. REFERENCES

[1] Miller, G. A. (1995). WordNet: A Lexical Database for English. Communications of the ACM, 38(11), 39-41. Available: https://dl.acm.org/doi/10.1145/219717.219748

[2] Salton, G., & McGill, M. J. (1983). Introduction to Modern Information Retrieval. McGraw-Hill. Available: https://archive.org/details/introductiontomo00salt

[3] Deerwester, S., Dumais, S. T., Furnas, G. W., Landauer, T. K., & Harshman, R. (1990). Indexing by Latent Semantic Analysis. Journal of the American Society for Information Science, 41(6), 391-407. Available: https://asistdl.onlinelibrary.wiley.com/doi/10.1002/(SICI)1097-4571(199009)41:6<391::AID-ASI1>3.0.CO;2-9

[4] Mikolov, T., Chen, K., Corrado, G., & Dean, J. (2013). Efficient Estimation of Word Representations in Vector Space. arXiv preprint arXiv:1301.3781. Available: https://arxiv.org/abs/1301.3781

[5] Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global Vectors for Word Representation. Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP), 1532–1543. Available: https://aclanthology.org/D14-1162/

[6] Devlin, J., Chang, M.-W., Lee, K., & Toutanova, K. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. Proceedings of NAACL-HLT, 4171–4186. Available: https://aclanthology.org/N19-1423/

[7] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., et al. (2020). Language Models Are Few-Shot Learners. Advances in Neural Information Processing Systems (NeurIPS). Available: https://arxiv.org/abs/2005.14165

[8] Samek, W., et al. (2021). "Explainable AI: Interpreting, Explaining and Visualizing Deep Learning." Available: https://www.researchgate.net/publication/335712512_Explainable_AI_Interpreting_Explaining_and_Visualizing_Deep_Learning